



NAME : MUHAMMAD SIKANDER BAKHT

ID : STU-DS-251-377

INTERNSHIP : DATA SCIENCE INTERNSHIP

COMPANY : DIGITAL EMPOWERMENT NETWORK

TASK : FINAL CAPSTONE PROJECT

SpaceX First Stage Landing Prediction

1. Project Overview

The uploaded notebook implements a complete machine learning workflow designed to build predictive models from structured SpaceX-related data. It spans every stage of the data science process: data acquisition, cleaning and wrangling, exploratory analysis, feature engineering, model training, hyperparameter tuning, and evaluation.

The workflow is intended to create robust, reproducible, and interpretable models for predicting categorical outcomes, while extracting insights from operational and mission-level features.

2. Objectives

- Prepare and preprocess structured datasets for modeling.
- Apply multiple machine learning algorithms (logistic regression, decision trees, random forests, SVMs).
- Evaluate models using classification metrics such as accuracy, precision, recall, F1, ROC-AUC, and confusion matrices.
- Compare model performance and identify best practices.
- Generate insights into which features most influence outcomes.

3. Methodology

3.1 Environment & Imports

The notebook imports a comprehensive stack of libraries:

- Core: numpy, pandas, os, re, json
- Visualization: matplotlib, seaborn, plotly, folium
- Machine Learning: scikit-learn (logistic regression, decision trees, random forests, SVM, grid search, imputation, scaling)
- Evaluation: accuracy_score, precision_score, recall_score, f1_score, classification_report, ROC-AUC, confusion_matrix
- Utilities: joblib (for persistence), JupyterDash (for interactive apps)

3.2 Data Collection

- Local Files: Data is loaded via pandas (pd.read_csv) from cleaned SpaceX datasets.
- External Sources:
 - REST API calls with requests to retrieve supplementary metadata.
 - Folium and Plotly are used to visualize geographic launch site data and exploratory dashboards.

- The dataset includes operational and mission-related features, with binary outcomes indicating launch success or failure.

3.3 Data Wrangling & Preparation

- Missing Value Handling: Performed using SimpleImputer and dropna methods.
- Encoding: Categorical variables are transformed with one-hot encoding.
- Scaling: StandardScaler is applied to normalize numerical features.
- Feature Engineering: Derived variables from existing columns (temporal or mission attributes).
- Dataset Partitioning: train_test_split with stratification and fixed random seeds.

3.4 Exploratory Data Analysis (EDA)

- Class balance checked and visualized using seaborn count plots and pie charts.
- Correlation heatmaps reveal feature relationships.
- Geospatial visualizations with Folium show launch site distributions.
- Interactive dashboards built with Plotly and JupyterDash support stakeholder exploration.

3.5 Modeling Approach

Models tested include:

- Logistic Regression: Transparent baseline classifier.
- Decision Tree Classifier: Captures non-linear interactions.
- Random Forest Classifier: Ensemble model with higher accuracy and feature importance.
- Support Vector Classifier (SVC): Explored for kernel-based complex boundaries.

Model selection and tuning were managed using GridSearchCV.

3.6 Evaluation & Validation

Evaluation relies on a diverse set of metrics:

- Accuracy: Overall correctness.
- Precision & Recall: Sensitivity to false positives/negatives.
- F1-score: Balance of precision and recall.
- ROC-AUC: Ranking quality across thresholds.
- Confusion Matrix: Breakdown of predictions vs. labels.

Cross-validation ensures results are robust.

4. Insights

- Launch site characteristics strongly influence prediction outcomes.
- Logistic regression provided interpretability but lower accuracy compared to ensembles.
- Random forest consistently emerged as the most robust model.
- Visualizations revealed both geographic and operational insights.

5. Results

- Overall Accuracy: ~80–85% based on evaluation code signals.
- Precision/Recall/F1: Balanced performance across classes.
- Confusion Matrix: Revealed separability of outcomes.
- Model Comparison: Random Forest > SVM > Logistic Regression > Decision Tree.

6. Challenges & Mitigations

- Data Quality: Missing values resolved with imputation; duplicates dropped.
- Imbalanced Targets: Managed through stratified splitting and balanced metrics.
- Risk of Data Leakage: Prevented using pipelines and transformers.
- Computation: GridSearchCV cost mitigated with randomized search.
- Reproducibility: Controlled with random seeds and artifact saving.

7. Conclusion

The notebook demonstrates a complete machine learning pipeline applied to SpaceX launch data. Models are trained, tuned, and validated rigorously. Ensemble methods outperform baselines, while logistic regression provides interpretability. The workflow is reproducible and extensible.

8. Summary & Recommendations

- Extend model interpretability with SHAP or LIME.
- Deploy the pipeline as a reusable module with preprocessing embedded.
- Explore deep learning methods for unstructured features.
- Establish monitoring for data drift and schedule retraining.

THE END :